

Day 19 – MovieLens Data Analysis with Pandas

In this notebook, I worked with the **MovieLens 20M dataset** — a popular real-world dataset for building and analyzing movie recommendation systems

Dataset Files Used:

- `movies.csv` → Contains movie metadata like title and genres
- `ratings.csv` → Contains user ratings for movies
- `tags.csv` → Contains user-generated tags associated with movies

All datasets were loaded using **Pandas**, and basic exploratory operations such as shape, column extraction, and cleaning were performed.

Objectives:

- Load and inspect multiple CSV files
- Clean and prepare the data
- Understand relationships between users, movies, and ratings

```
In [1]: import pandas as pd
```

```
In [2]: movies = pd.read_csv(r"C:\Users\Arman\Downloads\dataset\archive\movie.csv")
```

```
In [3]: print(type(movies))
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
In [4]: movies.head()
```

```
Out[4]:
```

	movielid	title	genres
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
1	2	Jumanji (1995)	Adventure Children Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama Romance
4	5	Father of the Bride Part II (1995)	Comedy

```
In [5]: movies.shape
```

```
Out[5]: (27278, 3)
```

```
In [6]: ratings = pd.read_csv(r"C:\Users\Arman\Downloads\dataset\archive\rating.csv")
```

```
In [7]: ratings.head()
```

```
Out[7]:
```

	userId	movieId	rating	timestamp
0	1	2	3.5	2005-04-02 23:53:47
1	1	29	3.5	2005-04-02 23:31:16
2	1	32	3.5	2005-04-02 23:33:39
3	1	47	3.5	2005-04-02 23:32:07
4	1	50	3.5	2005-04-02 23:29:40

```
In [8]: ratings.shape
```

```
Out[8]: (20000263, 4)
```

```
In [9]: tags = pd.read_csv(r"C:\Users\Arman\Downloads\dataset\archive\tag.csv")
```

```
In [10]: tags.head()
```

```
Out[10]:
```

	userId	movieId	tag	timestamp
0	18	4141	Mark Waters	2009-04-24 18:19:40
1	65	208	dark hero	2013-05-10 01:41:18
2	65	353	dark hero	2013-05-10 01:41:19
3	65	521	noir thriller	2013-05-10 01:39:43
4	65	592	dark hero	2013-05-10 01:41:18

```
In [11]: tags.shape
```

```
Out[11]: (465564, 4)
```

```
In [12]: tags.columns
```

```
Out[12]: Index(['userId', 'movieId', 'tag', 'timestamp'], dtype='object')
```

```
In [13]: ratings.columns
```

```
Out[13]: Index(['userId', 'movieId', 'rating', 'timestamp'], dtype='object')
```

```
In [14]: del ratings['timestamp']
del tags['timestamp']
```

```
In [15]: ratings
```

Out[15]:

	userId	movieId	rating
0	1	2	3.5
1	1	29	3.5
2	1	32	3.5
3	1	47	3.5
4	1	50	3.5
...
20000258	138493	68954	4.5
20000259	138493	69526	4.5
20000260	138493	69644	3.0
20000261	138493	70286	5.0
20000262	138493	71619	2.5

20000263 rows × 3 columns

In [16]: `tags`

Out[16]:

	userId	movieId	tag
0	18	4141	Mark Waters
1	65	208	dark hero
2	65	353	dark hero
3	65	521	noir thriller
4	65	592	dark hero
...
465559	138446	55999	dragged
465560	138446	55999	Jason Bateman
465561	138446	55999	quirky
465562	138446	55999	sad
465563	138472	923	rise to power

465564 rows × 3 columns

Data Structure

In [17]: `tags.head()`

Out[17]:

	userId	movieId	tag
0	18	4141	Mark Waters
1	65	208	dark hero
2	65	353	dark hero
3	65	521	noir thriller
4	65	592	dark hero

```
In [18]: tags.iloc[2]
```

```
Out[18]: userId      65
movieId     353
tag        dark hero
Name: 2, dtype: object
```

```
In [19]: row_0 = tags.iloc[0]
row_0
```

```
Out[19]: userId      18
movieId    4141
tag        Mark Waters
Name: 0, dtype: object
```

```
In [20]: type(row_0)
```

```
Out[20]: pandas.core.series.Series
```

```
In [21]: row_0.index
```

```
Out[21]: Index(['userId', 'movieId', 'tag'], dtype='object')
```

```
In [22]: print(row_0['userId'])
```

```
18
```

```
In [23]: 'rating' in row_0
```

```
Out[23]: False
```

```
In [24]: row_0.name
```

```
Out[24]: 0
```

```
In [25]: row_0 = row_0.rename('firstRow')
row_0
```

```
Out[25]: userId      18
movieId    4141
tag        Mark Waters
Name: firstRow, dtype: object
```

```
In [26]: row_0.name
```

```
Out[26]: 'firstRow'
```

DataFrames

```
In [27]: tags.head()
```

```
Out[27]:   userId  movieId       tag
0      18      4141  Mark Waters
1      65      208  dark hero
2      65      353  dark hero
3      65      521  noir thriller
4      65      592  dark hero
```

```
In [28]: tags.index
```

```
Out[28]: RangeIndex(start=0, stop=465564, step=1)
```

```
In [29]: tags.columns
```

```
Out[29]: Index(['userId', 'movieId', 'tag'], dtype='object')
```

```
In [30]: tags.iloc[[0,11,500]]
```

```
Out[30]:
```

	userId	movieId	tag
0	18	4141	Mark Waters
11	65	1783	noir thriller
500	342	55908	entirely dialogue

Descriptive Statistics

```
In [31]: ratings
```

```
Out[31]:
```

	userId	movieId	rating
0	1	2	3.5
1	1	29	3.5
2	1	32	3.5
3	1	47	3.5
4	1	50	3.5
...
20000258	138493	68954	4.5
20000259	138493	69526	4.5
20000260	138493	69644	3.0
20000261	138493	70286	5.0
20000262	138493	71619	2.5

20000263 rows × 3 columns

```
In [32]: ratings['rating'].describe()
```

```
Out[32]:
```

count	2.000026e+07
mean	3.525529e+00
std	1.051989e+00
min	5.000000e-01
25%	3.000000e+00
50%	3.500000e+00
75%	4.000000e+00
max	5.000000e+00
Name:	rating, dtype: float64

```
In [33]: ratings.describe()
```

```
Out[33]:
```

	userId	movieId	rating
count	2.000026e+07	2.000026e+07	2.000026e+07
mean	6.904587e+04	9.041567e+03	3.525529e+00
std	4.003863e+04	1.978948e+04	1.051989e+00
min	1.000000e+00	1.000000e+00	5.000000e-01
25%	3.439500e+04	9.020000e+02	3.000000e+00
50%	6.914100e+04	2.167000e+03	3.500000e+00
75%	1.036370e+05	4.770000e+03	4.000000e+00
max	1.384930e+05	1.312620e+05	5.000000e+00

```
In [34]: ratings['rating'].mean()
```

```
Out[34]: np.float64(3.5255285642993797)
```

```
In [35]: ratings.mean()
```

```
Out[35]: userId      69045.872583
movieId     9041.567330
rating       3.525529
dtype: float64
```

```
In [36]: ratings['rating'].min()
```

```
Out[36]: 0.5
```

```
In [37]: ratings['rating'].max()
```

```
Out[37]: 5.0
```

```
In [38]: ratings['rating'].std()
```

```
Out[38]: 1.051988919275684
```

```
In [39]: ratings['rating'].mode()
```

```
Out[39]: 0    4.0
Name: rating, dtype: float64
```

```
In [40]: ratings.corr()
```

```
Out[40]:
```

	userId	movieId	rating
userId	1.000000	-0.000850	0.001175
movieId	-0.000850	1.000000	0.002606
rating	0.001175	0.002606	1.000000

```
In [41]: filter1 = ratings['rating'] > 10
print(filter1)
```

```
0      False
1      False
2      False
3      False
4      False
...
20000258  False
20000259  False
20000260  False
20000261  False
20000262  False
Name: rating, Length: 20000263, dtype: bool
```

```
In [42]: filter1.any()
```

```
Out[42]: np.False_
```

```
In [43]: filter2 = ratings['rating'] > 0
print(filter2)
```

```
0      True
1      True
2      True
3      True
4      True
...
20000258  True
20000259  True
20000260  True
20000261  True
20000262  True
Name: rating, Length: 20000263, dtype: bool
```

```
In [44]: print(filter2.all())
```

```
True
```

Data Cleaning: Handling Missing Data

```
In [45]: movies
```

	moviedb_id	title	genres
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
1	2	Jumanji (1995)	Adventure Children Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama Romance
4	5	Father of the Bride Part II (1995)	Comedy
...
27273	131254	Kein Bund für's Leben (2007)	Comedy
27274	131256	Feuer, Eis & Dosenbier (2002)	Comedy
27275	131258	The Pirates (2014)	Adventure
27276	131260	Rentun Ruusu (2001)	(no genres listed)
27277	131262	Innocence (2014)	Adventure Fantasy Horror

```
27278 rows × 3 columns
```

```
In [46]: movies.shape
```

```
Out[46]: (27278, 3)
```

```
In [47]: movies.isnull().any().any()
```

```
Out[47]: np.False_
```

```
In [48]: ratings.shape
```

```
Out[48]: (20000263, 3)
```

```
In [49]: ratings.isnull().any().any()
```

```
Out[49]: np.False_
```

```
In [50]: tags.shape
```

```
Out[50]: (465564, 3)
```

```
In [51]: tags.isnull().any().any()
```

```
Out[51]: np.True_
```

```
In [52]: tags = tags.dropna()
```

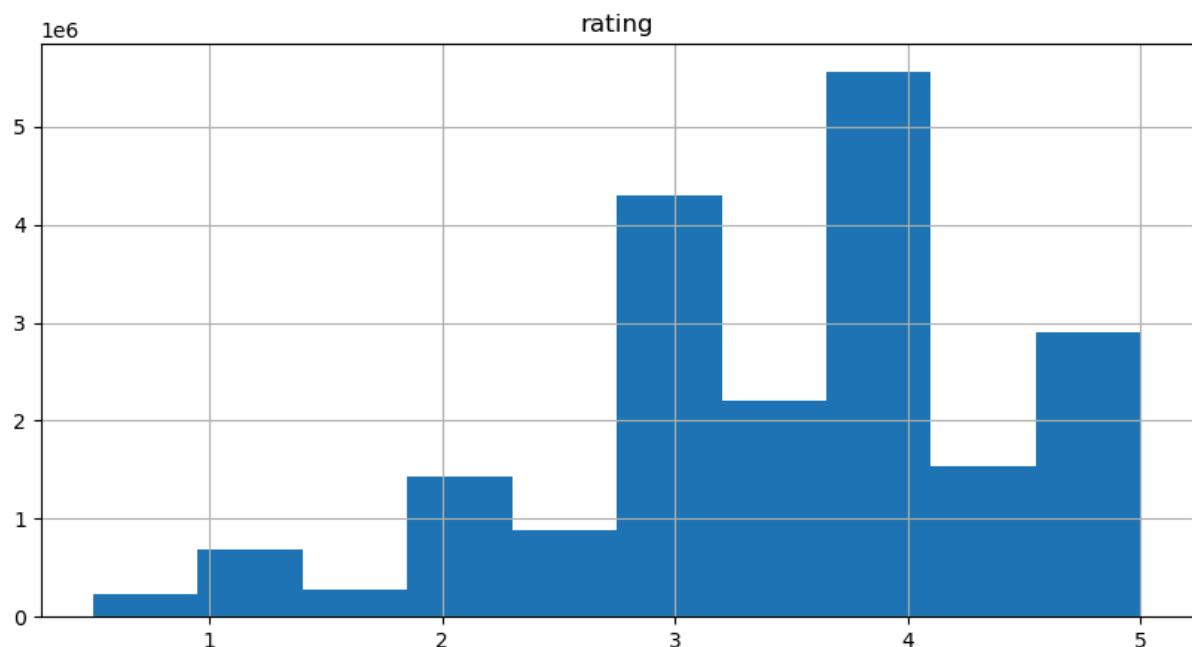
```
In [53]: tags.isnull().any().any()
```

```
Out[53]: np.False_
```

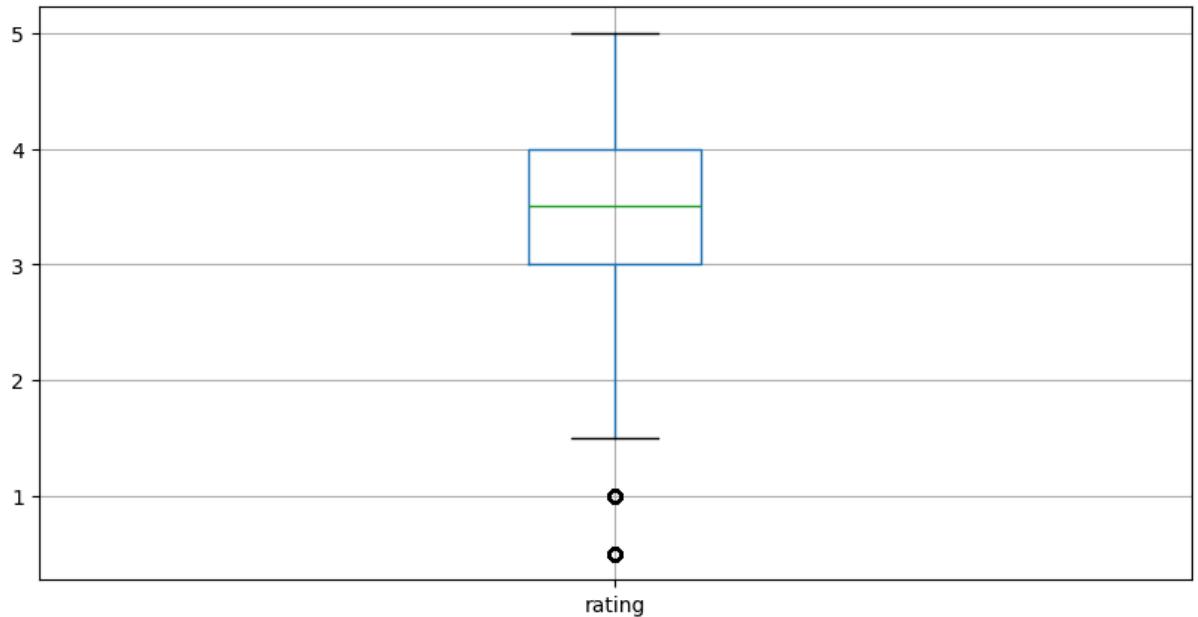
Data Visualization

```
In [54]: import matplotlib.pyplot as plt  
%matplotlib inline
```

```
In [55]: ratings.hist(column='rating', figsize=(10,5))  
plt.show()
```



```
In [56]: ratings.boxplot(column='rating', figsize=(10,5))  
plt.show()
```



Slicing Out Columns

```
In [57]: tags['tag'].head()
```

```
Out[57]: 0      Mark Waters
1      dark hero
2      dark hero
3    noir thriller
4      dark hero
Name: tag, dtype: object
```

```
In [58]: movies[['title','genres']].head()
```

```
Out[58]:      title           genres
0  Toy Story (1995)  Adventure|Animation|Children|Comedy|Fantasy
1  Jumanji (1995)        Adventure|Children|Fantasy
2  Grumpier Old Men (1995)          Comedy|Romance
3  Waiting to Exhale (1995)  Comedy|Drama|Romance
4  Father of the Bride Part II (1995)          Comedy
```

```
In [59]: ratings[-10:]
```

```
Out[59]:
```

	userId	movieId	rating
20000253	138493	60816	4.5
20000254	138493	61160	4.0
20000255	138493	65682	4.5
20000256	138493	66762	4.5
20000257	138493	68319	4.5
20000258	138493	68954	4.5
20000259	138493	69526	4.5
20000260	138493	69644	3.0
20000261	138493	70286	5.0
20000262	138493	71619	2.5

```
In [60]: tag_counts = tags['tag'].value_counts()  
tag_counts[-10:]
```

```
Out[60]: tag
```

```
Hell naw          1  
This is my happy face    1  
I heel toe on Uday's house 1  
Why?             1  
Bobo              1  
Diamond Dallas Page    1  
I'm Devon Butler!      1  
No arguement        1  
Really Bad          1  
Botox              1  
Name: count, dtype: int64
```

```
In [61]: tag_counts[:10].plot(kind='bar', figsize=(10,5))  
plt.show()
```

